

# Cache Lab

**Instructors:**

Sungyong Ahn

# What is Cache LAB

- This lab will help you understand the impact that cache memories can have on the performance of your C programs.
- Part A: Writing a Cache Simulator
  - Write a small C program (about 200-300 lines) that simulates the behavior of a cache memory
- Part B: Optimizing Matrix Transpose
  - Optimize a small matrix transpose function, with the goal of minimizing the number of cache misses

# Handout Instructions

- Lab materials are contained in a Unix tar file called *cachelab-handout.tar*
  - You can download from **PLATO**
- You have to extract the tar file in Linux
  - `$> tar xvf cachelab-handout.tar`
- You will be modifying two files: `csim.c` and `trans.c`.
- **WARNING:** Do not let the Windows WinZip program open up your .tar file

# Reference Trace Files

- The `traces` subdirectory of the handout directory contains a collection of reference trace files generated by a Linux program called `valgrind`
- `[space]operation address, size`
- 64-bit hexadecimal memory address

```
I 0400d7d4, 8  
M 0421c7f0, 4  
L 04f6b868, 8  
S 7ff0005c8, 8
```

I: instruction load

L: data load

S: data store

M: data modify

NOTE: There is never a space before each “I”

# Part A: Writing a Cache Simulator

- Write a cache simulator in `csim.c`
  - Input: a `valgrind` memory trace
  - Simulates the hit/miss behavior of a cache memory
  - Outputs: the total number of hits, misses, and evictions.

# Part A: Writing a Cache Simulator

- Reference cache simulator (`csim-ref`)
  - Usage: `./csim-ref [-hvv] -s <s> -E <E> -b <b> -t <tracefile>`
  - `-h`: Optional help flag that prints usage info
  - `-v`: Optional verbose flag that displays trace info
  - `-s <s>`: Number of set index bits ( $S = 2^s$  is the number of sets)
  - `-E <E>`: Associativity (number of lines per set)
  - `-b <b>`: Number of block bits ( $B = 2^b$  is the block size)
  - `-t <tracefile>`: Name of the `valgrind` trace to replay

```
linux> ./csim-ref -s 4 -E 1 -b 4 -t traces/yi.trace
hits:4 misses:5 evictions:3
```

# Part A: Writing a Cache Simulator

## ■ Programming Rules for Part A

- Include your name and loginID in the header comment for `csim.c`.
- Your `csim.c` file must compile without warnings in order to receive credit
- Ignore all instruction cache accesses (lines starting with “I”)
- The data modify operation (M) is treated as a load followed by a store to the same address.
- Assume that memory accesses are aligned properly
  - you can ignore the request sizes in the `valgrind` traces

```
/* Kildong Hong, 201912345 */  
#include <getopt.h>  
#include <stdlib.h>  
#include <unistd.h>  
#include <stdio.h>
```

## Part B: Optimizing Matrix Transpose

- Write a transpose function, called `transpose_submit` in `trans.c` that causes as few cache misses as possible
- The transpose of  $A$ , denoted  $A^T$ , is a matrix such that  $A_{ij} = A^T_{ji}$
- Example transpose function

```
char trans_desc[] = "Simple row-wise scan transpose";  
void trans(int M, int N, int A[N][M], int B[M][N])
```

- Your job to write a similar function, called `transpose_submit`

```
char transpose_submit_desc[] = "Transpose submission";  
void transpose_submit(int M, int N, int A[N][M], int B[M][N]);
```

**Warning:** Do not change the description string ("Transpose submission")



# Part B: Optimizing Matrix Transpose

## ■ Programming Rules for Part B

- Include your name and loginID in the header comment for `trans.c`.
- Your code in `trans.c` must compile without warnings to receive credit
- You are allowed to define at most **12 local variables** of type **int** per transpose function
- Your transpose function may not use recursion
- Your transpose function may not modify array A. You may, however, do whatever you want with the contents of array B.
- You are NOT allowed to define any arrays in your code or to use any variant of malloc.

# Evaluation (Part A: 27)

- Autograding program, called `test-csim`

```
linux> make
linux> ./test-csim
```

Your simulator		Reference simulator						
Points	(s,E,b)	Hits	Misses	Evicts	Hits	Misses	Evicts	
0	(1,1,1)	0	0	0	9	8	6	traces/yi2.trace
0	(4,2,4)	0	0	0	4	5	2	traces/yi.trace
0	(2,1,4)	0	0	0	2	3	1	traces/dave.trace
0	(2,1,3)	0	0	0	167	71	67	traces/trans.trace
0	(2,2,3)	0	0	0	201	37	29	traces/trans.trace
0	(2,4,3)	0	0	0	212	26	10	traces/trans.trace
1	(5,1,5)	0	0	0	231	7	0	traces/trans.trace
0	(5,1,5)	0	0	0	265189	21775	21743	traces/long.trace

```
1

TEST_CSIM_RESULTS=1
```

# Evaluation (Part B: 26)

- Autograding program, called `test-trans.c`
- You can register up to 100 versions of the transpose function in your `trans.c` file.

```
/*
 * You can define additional transpose functions below. We've defined
 * a simple one below to help you get started.
 */
char transpose_test_desc[] = "Transpose test";
void transpose_test(int M, int N, int A[N][M], int B[M][N])
{
    /* Insert matrix transpose code you want to test */
}

void registerFunctions()
{
    ...
    /* Register any additional transpose functions */
    registerTransFunction(transpose_test, transpose_test_desc);
    ...
}
```

# Evaluation (Part B: 26)

- Autograding program, called `test-trans.c`
- You can register up to 100 versions of the transpose function in your `trans.c` file.
- You must implement **two versions of the transpose function.**

```
linux> make
linux> ./test-trans -M 32 -N 32
Function 0 (4 total)
Step 1: Validating and generating memory traces
Step 2: Evaluating performance (s=5, E=1, b=5)
func 0 (Transpose submission): hits:1766, misses:287, evictions:255
Function 1 (4 total)
Step 1: Validating and generating memory traces
Step 2: Evaluating performance (s=5, E=1, b=5)
func 1 (Simple row-wise scan transpose): hits:870, misses:1183, evictions:1151

Summary for official submission (func 0): correctness=1 misses=287
```

# Evaluation (Putting it all Together)

- *driver program*, called `./driver.py`, that performs a complete evaluation of your simulator and transpose code.

# Handing in Your Work

- Each time you type `make` in the `cachelab-handout` directory, the Makefile creates a tarball, called `userid-handin.tar`, that contains your current `csim.c` and `trans.c` files.
- You need to upload your `userid-handin.tar` file, to **PLATO**.
- Note that you must implement two transpose function including `transpose_submit` function

# Submission

- Due to **6/14 (Mon.)** 23:59
  - 하루 딜레이 시 만점기준 25% 감점
- `userid-handin.tar` 파일 PLATO에 제출
- Please Read the Writup.